

Cadence Sample Document

Introduction

Backpropagation computes gradients via the chain rule. It is the workhorse of modern deep learning. The algorithm propagates error signals backward through each layer.

Dense passages of technical prose contain longer polysyllabic terminology that a thoughtful reader naturally slows down for. Cadence mirrors that instinct automatically.

Optimization

Gradient descent nudges each weight against its gradient. A learning rate that is too large will diverge. A rate that is too small will crawl.

Momentum accumulates a running average of past gradients. This helps the optimizer glide through shallow valleys.

Regularization

Weight decay penalizes large parameters to curb overfitting. Dropout randomly removes units during training. Both encourage models that generalize.

Early stopping halts training once validation loss plateaus. It is simple and surprisingly effective.